

Ahmed Refaay

Nov. 5th, 2024

Project 1 Report: Transfer Learning with MobileNetV2 for CIFAR-10

1. Introduction

This project focuses on applying transfer learning techniques to the CIFAR-10 dataset using MobileNetV2, a lightweight and efficient deep learning model ideal for resource-constrained environments. The project's goals align with companies' emphasis on optimization and practical deployment in real-world scenarios.

2. Project Overview

- **Objective:** To develop a high-accuracy, resource-efficient model for classifying CIFAR-10 images, demonstrating skills in model customization, optimization, and practical deployment.
 - **Model Choice:** MobileNetV2 was selected for its balance between accuracy and efficiency, which is essential for low-resource setups used in edge applications.
-

3. Methodology

1. Data Loading and Preprocessing:

- **Dataset:** CIFAR-10, consisting of 60,000 32x32 images across 10 classes.
- **Resizing:** Images were resized to 96x96 to be compatible with MobileNetV2.
- **Data Augmentation:** Techniques such as random rotations, horizontal flipping, and scaling were used to enhance model generalization.

2. Model Setup and Training:

- **Base Model:** Loaded MobileNetV2 without the top layers, freezing the pre-trained weights to retain learned features.
- **Custom Layers:** Added a fully connected layer setup to tailor the model for CIFAR-10 classification.
- **Compilation:** Used the Adam optimizer, categorical cross-entropy loss, and set metrics to track accuracy.
- **Training and Fine-Tuning:**
 - Initially trained with frozen layers to establish a baseline.
 - Fine-tuned by unfreezing a portion of the MobileNetV2 layers to further improve performance.
- **Outcome:** Achieved a testing accuracy of 91.89%, reflecting a robust model suitable for efficient inference.

3. Inference and Evaluation:

- **Testing:** Conducted inference on new images from the internet, confirming accurate classification.
 - **Performance Metrics:** Monitored training and validation loss curves, along with accuracy trends.
-

4. Challenges and Solutions

- **Memory Management:** Encountered issues with memory exhaustion on GPU/CPU. Addressed this by using data augmentation in batches and optimizing TensorFlow settings.
 - **Resource Utilization:** Explored strategies to balance CPU and GPU load, ensuring smoother training and inference processes.
-

5. Future Directions

1. Optimization:

- Explore using OpenVINO or TensorRT for performance improvements on edge devices.
- Consider model quantization or pruning to reduce the model's size and increase efficiency.

2. Deployment:

- Package the model using Docker for easy deployment and integration.
 - Conduct tests in simulated production environments to ensure stability and performance.
-

6. Conclusion

This project demonstrates a thorough understanding of transfer learning and model optimization for low-resource environments. By achieving high accuracy while maintaining efficiency, the project showcases practical skills relevant to companies' focus on high-performance computer vision solutions. The next steps include further optimization and real-world deployment practices.

Attachments:

- Full implementation code
- Training and evaluation plots (log files)
- Documentation of the deployment setup (to be prepared)